

Master Degree in Artificial Intelligence

# **Boids Simulation: Parallel Programming Optimization**


**Parallel Programming for Machine Learning 2025**

Gianni Moretti

# Introduction

# What are Boids?

- Introduced by Craig Reynolds in 1986
- Simulates flocking behavior of birds, fish, or other organisms



# The Three Fundamental Rules

- **Alignment:** Each boid tries to match its velocity (direction and speed) to the average velocity of its local neighbors. This causes the group to move in a similar direction, promoting coordinated movement.
- **Cohesion:** Each boid steers towards the average position (center of mass) of its nearby flockmates. This keeps the group together and prevents individuals from straying too far from the flock.
- **Separation:** Each boid actively avoids crowding by steering away from neighbors that are too close. This prevents collisions and maintains a comfortable distance between individuals.

Each boid follows these rules independently, creating realistic flocking patterns without central control.

# Project Goals

- Implement efficient Boids simulation in C++
- Compare sequential vs parallel implementations
- Analyze different data layouts: Array of Structures (AoS) vs Structure of Arrays (SoA)
- Evaluate performance scalability with OpenMP
- Utilize SIMD optimizations where possible

# Implementation Overview

## Technologies Used:

- **Language:** C++ for performance
- **Parallelization:** OpenMP for multi-core optimization
- **Visualization:** SFML (Simple and Fast Multimedia Library)
- **Build System:** CMake for cross-platform compilation

# Implementation Details

# Data Layout Comparison: AoS vs SoA

## Array of Structures (AoS)

### Traditional object-oriented approach

```
1 struct BoidData {  
2     sf::Vector2f position;  
3     sf::Vector2f velocity;  
4     float biasval;  
5     int scoutGroup;  
6 };
```

- Easy to understand and maintain
- Less cache-friendly for bulk operations
- Memory accesses can be scattered



# Data Layout Comparison: AoS vs SoA

## Structure of Arrays (SoA)

### Performance-optimized layout

```
1 struct BoidDataList {  
2     float* xPos;  
3     float* yPos;  
4     float* xVelocity;  
5     float* yVelocity;  
6     float* biasvals;  
7     int* scoutGroup;  
8     int numBoid;  
9 };
```

- Better cache utilization
- Enables SIMD vectorization
- Contiguous memory access patterns

# Sequential Implementation

## Main Algorithm Structure:

```
1 for (int i = 0; i < num_boids; i++) {  
2     // Compute alignment, cohesion, separation  
3     for (int j = 0; j < num_boids; j++) {  
4         if (i != j && distance < perception_radius) {  
5             // Apply boid interaction rules  
6         }  
7     }  
8     // Update position and velocity  
9 }
```

## Characteristics:

- $O(N^2)$  computational complexity
- Perfect candidate for parallelization

# Parallel Implementation with OpenMP

## Key Parallelization Strategy:

```
1 #pragma omp parallel for schedule(static)
2 for (int i = 0; i < N; i++) {
3     // Each thread processes different boids
4     // Thread-safe access to read-only data
5     // Write to separate temporary arrays
6 }
7
8 #pragma omp simd
9 for (int i = 0; i < N; i++) {
10     // Vectorized copy back to main arrays
11     boidDataList.xPos[i] = new_xPos[i];
12     boidDataList.yPos[i] = new_yPos[i];
13     // ...
14 }
```

# Parallel Implementation with OpenMP

## Thread Safety Approach:

- Read-only access to current state
- Write results to temporary arrays
- Vectorized copy-back operation
- No race conditions or data dependencies

## Compiler Vectorization Enablers:

- Contiguous memory access (SoA layout)
- `#pragma omp simd` directives
- Multiple data elements processed per instruction

# Experimental Setup

# Benchmark Methodology

## Testing Environment:

- Multi-core CPU system (12 cores)
- Identical simulation parameters across all tests
- Multiple runs per configuration for statistical accuracy (30 runs)
- Isolated measurement of core algorithm (no rendering overhead)

## Test Parameters:

- **Boid populations:** 1,000 to 32,000 agents
- **Thread configurations:** 1 to 12 threads
- **Implementations:** Sequential AoS vs Sequential SoA vs Parallel SoA

# Performance Metrics

## Primary Measurements:

- Execution time per simulation step
- Speedup vs sequential baseline
- Scaling efficiency with thread count

## Analysis Focus:

- Impact of data layout optimization
- Parallel scalability characteristics
- Optimal thread configuration
- Performance vs problem size relationship

# Computational Complexity Analysis

## Algorithm Characteristics:

- **Time Complexity:**  $O(N^2)$  per simulation step
- **Space Complexity:**  $O(N)$  for boid storage
- **Parallel Potential:** Embarrassingly parallel outer loop



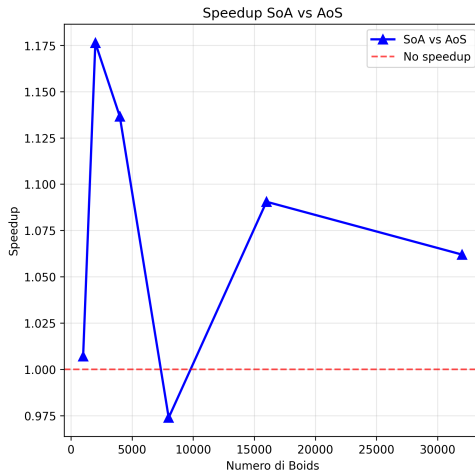
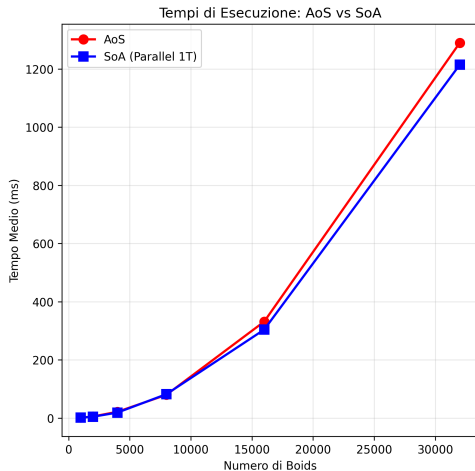
# **Performance Results**

# Data Layout Impact: AoS vs SoA

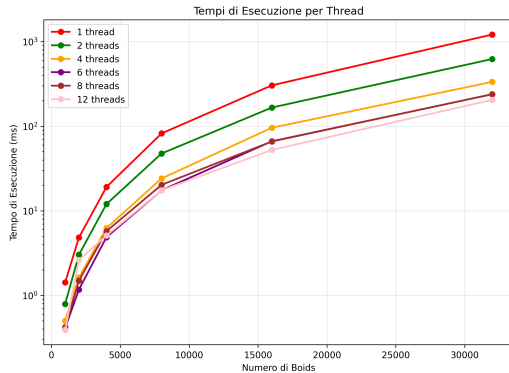
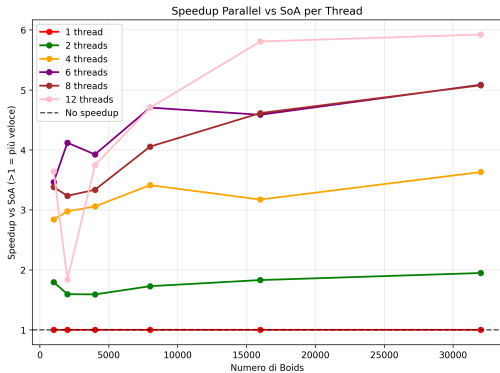
## Key Findings:

- SoA provides modest but consistent performance improvement
- 6.2% improvement for 32,000 boids (1,290.09ms vs 1,214.82ms)
- Benefits become more pronounced with larger datasets
- Foundation for effective parallelization

# Data Layout Impact: AoS vs SoA



# Parallel Scaling Performance



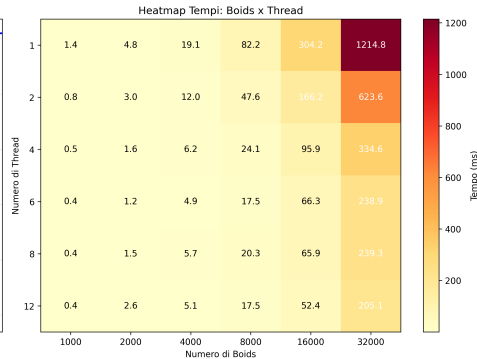
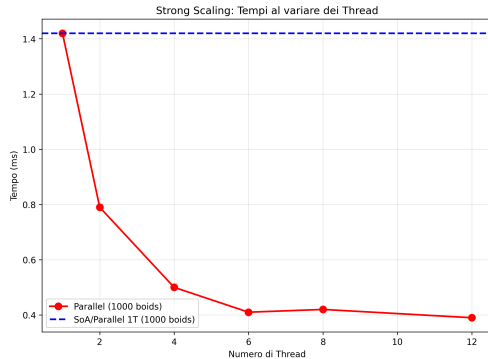
# Parallel Scaling Performance

## Speedup Results:

- Near-linear scaling up to 6 threads
- Diminishing returns beyond 8-12 threads due to overhead

<b>Dataset Size</b>	<b>Sequential (ms)</b>	<b>Best Parallel (ms)</b>	<b>Speedup</b>
1,000 boids	1.42	0.39	3.6x
4,000 boids	22.69	4.85	4.7x
8,000 boids	82.23	17.48	4.7x
16,000 boids	367.89	74.23	5.0x
32,000 boids	1,214.82	205.12	5.9x

# Parallel Scaling Performance



# Conclusions

## Project Achievements:

- Successfully implemented and optimized Boids simulation
- Demonstrated significant performance improvements through parallelization
- Validated importance of data structure optimization
- Achieved up to 5.9x speedup on multi-core systems

# **Thanks for your attention**

**Gianni Moretti**